// FIELD BRIEF

Smaller prompts, same answers

418 to 233 tokens per request

A plain request often carries more context than the model needs. This brief lists the habits that already help, five practices to add, and a two-line way to measure the result without whole-file dumps or long tool descriptions.

ALREADY IN PLAY

Three habits the team uses today.

- Short system notes
- Cached shared context
- Trimmed chat history

FIVE PRACTICES TO ADD

Each one is small and easy to check.

- Send diffs, not files
- Name only the tools in use
- Summarise old turns
- Reuse stable prefixes
- Cap the reply length

WHAT ONE REQUEST LOADS

41%
system notes

27%
tool text

22%
history

10%
the question

HOW TO MEASURE

1. Log the prompt size
2. Change one habit *same every time*
3. Run the same ten tasks
4. Compare size and quality

WHERE WE GO FROM HERE

Start with the two cheapest practices this week, then repeat the measurement and keep only the changes that hold quality steady.